

Maintaining Applications with Deployments



Anthony E. Nocentino

ENTERPRISE ARCHITECT @ CENTINO SYSTEMS

@nocentino www.centinosystems.com

Course Overview



Using Controllers to Deploy Applications and Deployment Basics

Maintaining Applications with Deployments

Deploying and Maintaining Applications with DaemonSets and Jobs

Overview

Configuring and Managing Application State with Deployments

- Updating Deployments
- Controlling Rollouts
- Scaling Applications


Managing Application State with Deployments



Creating




Updating



Scaling

Updating a Deployment



Rolling out a new
container image

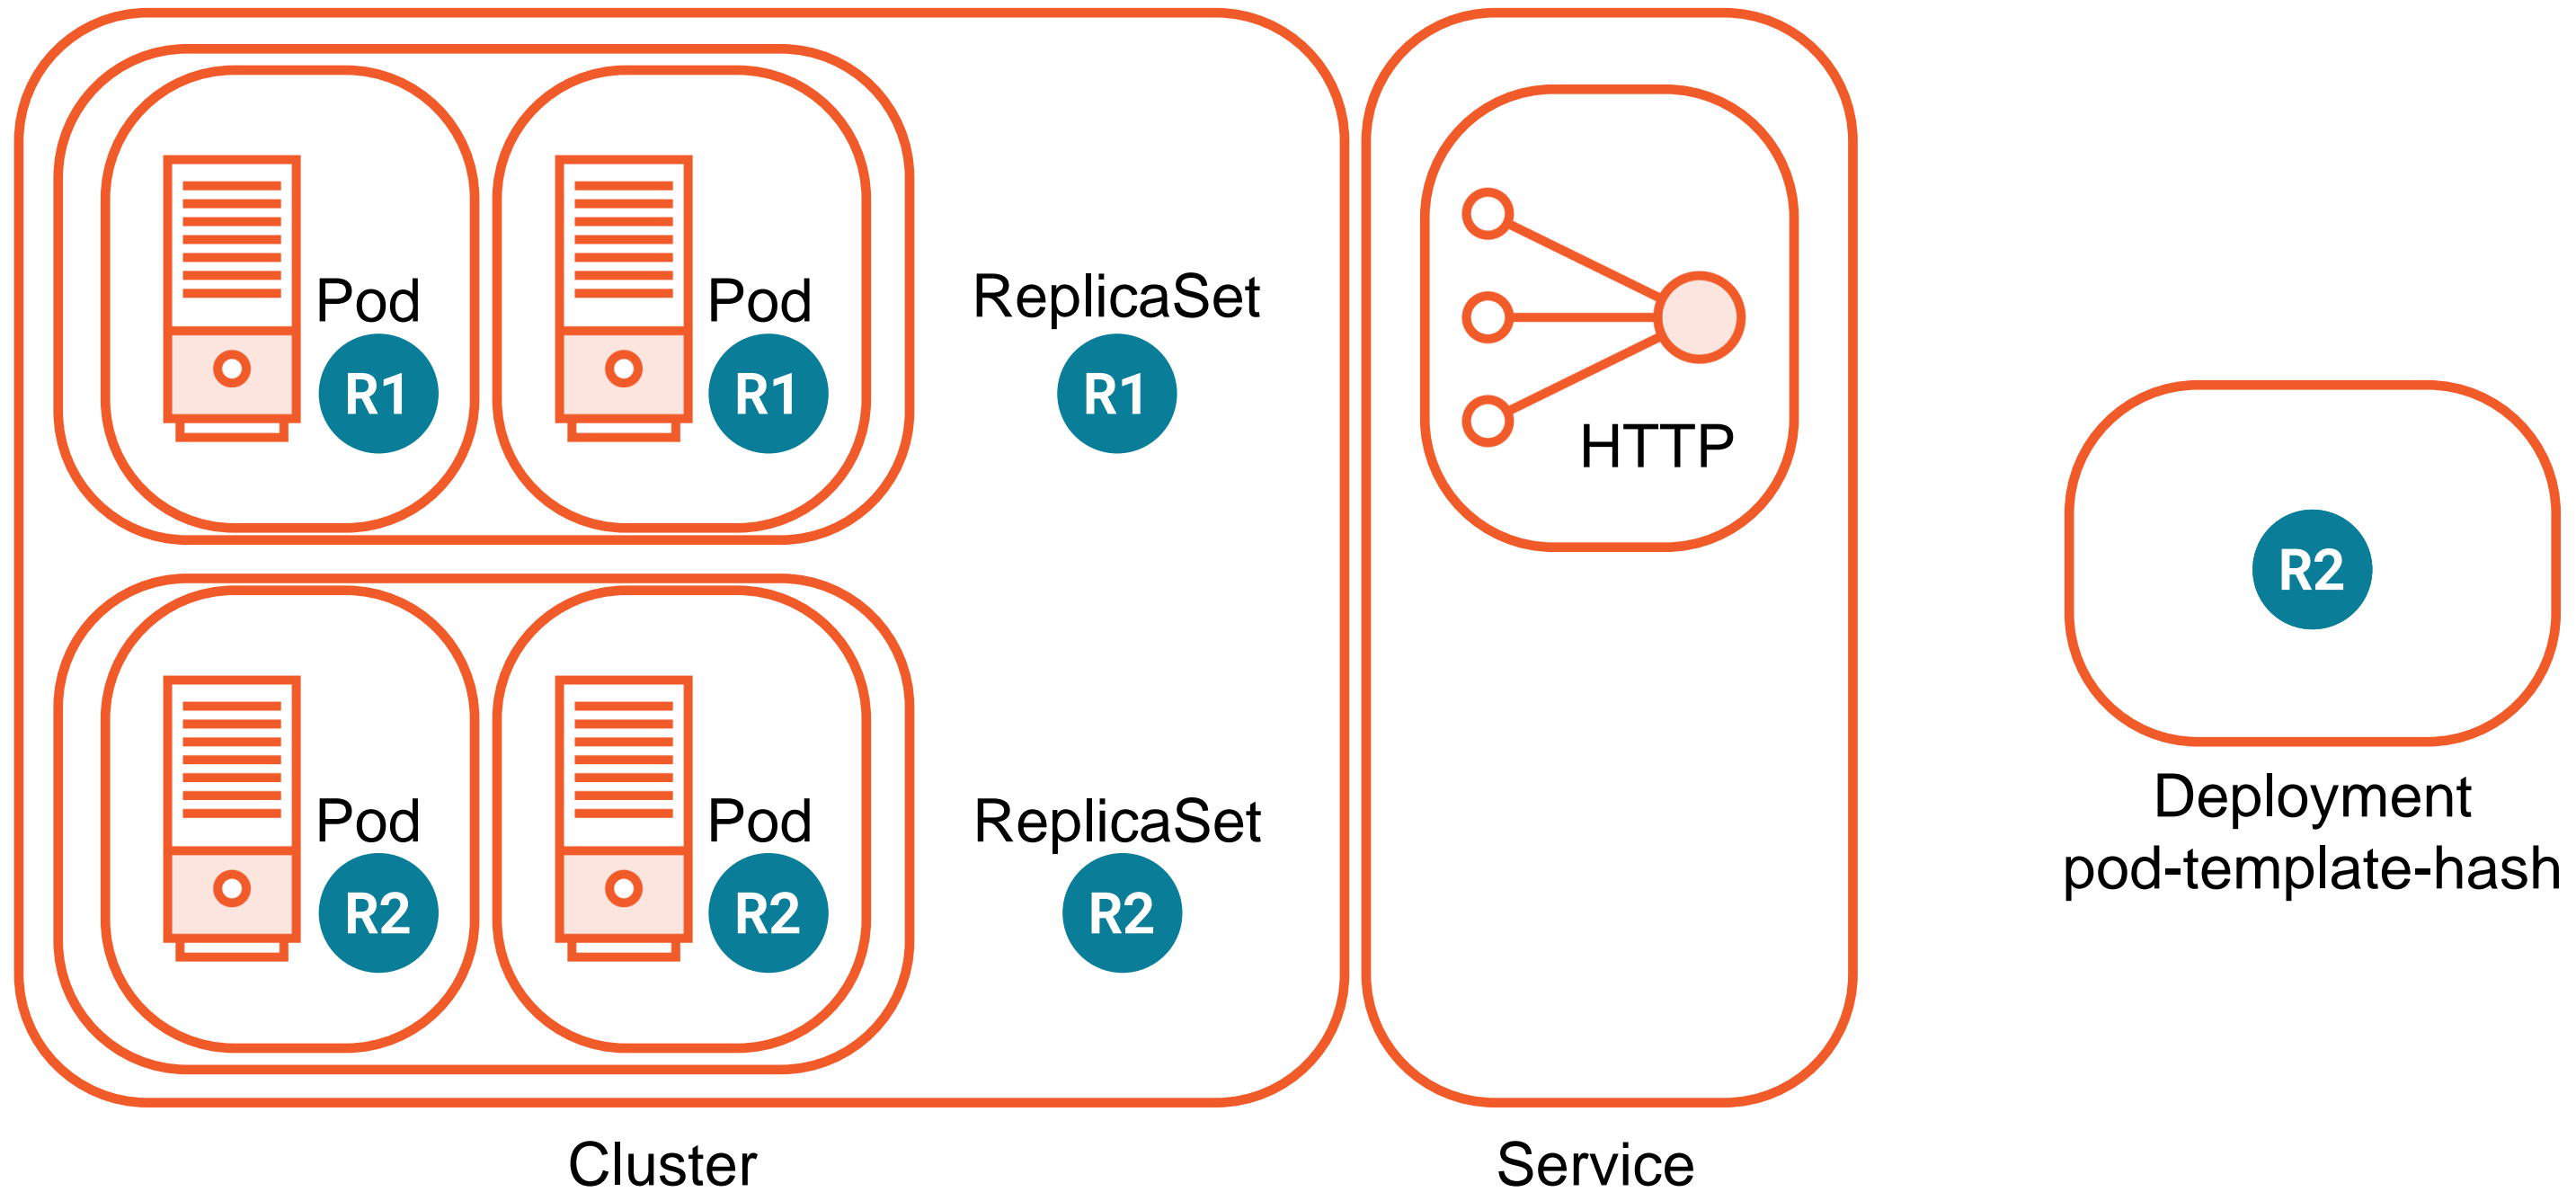


Triggered by changing
the Pod Template



Other fields can be
changed without
triggering an update

Controller Operations - Deployment Updates



Updating a Deployment Object

```
kubectl set image deployment hello-world hello-world=hello-app:2.0
```

```
kubectl set image deployment hello-world hello-world=hello-app:2.0 --record
```

```
kubectl edit deployment hello-world
```

```
kubectl apply -f hello-world-deployment.yaml --record
```

Checking Deployment Status



```
kubectl rollout status deployment [name]
```

```
kubectl describe deployment [name]
```

Deployment Status

Complete - all update work is finished

Progressing - update in flight


Failed - update could not complete

Demo

Updating a Deployment

Checking Deployment Rollout Status

Using Deployments to Change State



Control rollouts of a new version of your application


Update Strategy

Pause to make corrections

Rollback to an earlier version

Restart a Deployment

Controlling Rollouts With Update Strategy



Controls Pods rollout

RollingUpdate (Default)

A new ReplicaSet starts scaling up and the old ReplicaSet starts scaling down

Recreate

Terminates all Pods in the current ReplicaSet set prior to scaling up the new ReplicaSet

Used when applications don't support running different versions concurrently

Controlling the RollingUpdate Strategy

`maxUnavailable`

Ensures only a certain number of Pods are unavailable being updated

`maxSurge`

Ensure that only a certain number of Pods are created above the desired number of Pods

Successfully Controlling Deployment Rollouts



Update Strategy in a Deployment Spec


Readiness Probes in your Pod Template Spec

Update Strategy

```
apiVersion: apps/v1
kind: Deployment
...
spec:
  replicas: 20
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 20%
      maxSurge: 5
  ...
```

```
    template:
      ...
      spec:
        containers:
          ...
          readinessProbe:
            httpGet:
              path: /index.html
              port: 8080
              initialDelaySeconds: 10
              periodSeconds: 10
```

Pausing and Resuming a Deployment



Changes to the Deployment while paused are not rolled out

Batch changes together, then resume the rollout

The current state of the Deployment is maintained until it's resumed

Starts up a new ReplicaSet with the new changes

```
kubectl rollout pause deployment \
my-deployment
```

```
kubectl rollout resume deployment \
my-deployment
```

Rolling Back a Deployment



Rollout history

CHANGE-CAUSE Annotation Deployment

Revision History

`revisionHistoryLimit` defaults to 10

Number of `ReplicaSets` retained in history

Used for rolling back

Can be set to 0 for immediate cleanup

Rolling Back a Deployment (con't)




```
kubectl rollout history deployment \
hello-world
```

```
kubectl rollout history deployment \
hello-world --revision=1
```

```
kubectl rollout undo deployment
hello-world
```

```
kubectl rollout undo deployment \
hello-world --to-revision=1
```

Restarting a Deployment



Effectively restarts all the Pods

But no Pod is ever “recreated”

New ReplicaSet with the same Pod Spec

Uses Deployment’s Update Strategy

RollingUpdate

Recreate

```
kubectl rollout restart \  
deployment hello-world
```

Demo


Rolling Back a Deployment

Controlling the rate of a Rollout

Using Readiness Probes to Control Rollout

Restarting a Deployment

Scaling Deployments



Manual

```
kubectl scale deployment hello-world --replicas=10  
kubectl apply -f deployment.yaml
```



Horizontal Pod Autoscaler

Demo

Scaling a Deployment


Deployment Tips



Control your rollouts with an Update Strategy appropriate for your application



Use Readiness Probes for your application



Use the `--record` option to leave a trail of your work for others

Review

Configuring and Managing Application State with Deployments

- Using Deployments to Change State
- Controlling Rollouts
- Scaling Applications

What's Next!

Deploying and Maintaining Applications with DaemonSets and Jobs